Quiz 4 Submission

Total points   4/4   ?

A score of 3/4 or 4/4 is required to be considered to have "passed" a quiz. Please do not resubmit a quiz if you obtain a score of 3/4. You don't receive a final grade at the end of the course, so it will have no bearing on your certificate!

Your quiz will be graded and returned to you within a few minutes in most cases. However, it may take up to three weeks for your work to be imported into your Gradebook. Do be patient, please!

Quizzes (which are submitted via Google Forms and not submit50) will not show up as submitted in your Gradebook, until the scores have been imported, and even then will only show up if you have received a passing score.

Email *

justinattw@icloud.com

Name *

Justin Wong

edX Username *

justinattw

GitHub Username *

justinattw

City, State, Country

✓   Categorize the following as supervised learning, reinforcement   *1/1
learning, unsupervised learning, or not machine learning: A social
network's AI uses existing tagged photos of people to identify when
those people appear in new photos.

○   Not an example of machine learning

○   Reinforcement learning

◉   Supervised learning   ✓

○   Unsupervised learning

✓   Imagine a regression AI that makes the following predictions for the   *1/1
following 5 data points. What is the total L2 loss across all of these data
points (i.e., the sum of all the individual L2 losses for each data point)?

For data point 1, the true output is 2 and the AI predicted 4. For data point 2, the true output is 4 and
the AI predicted 5. For data point 3, the true output is 4 and the AI predicted 3. For data point 4, the
true output is 5 and the AI predicted 2. For data point 5, the true output is 6 and the AI predicted 5.

○   0

○   4

○   5

○   8

◉   16   ✓

○   19

○   21

○   64

✓   If Hypothesis 1 has a lower L1 loss and a lower L2 loss than Hypothesis 2   *1/1
on a set of training data, why might Hypothesis 2 still be a preferable
hypothesis?

○   Hypothesis 1 might be the result of regularization.

○   Hypothesis 1 might be the result of cross-validation.

○   Hypothesis 1 might be the result of loss.

◉   Hypothesis 1 might be the result of overfitting.   ✓

○   Hypothesis 1 might be the result of regression.

✓   In the ε-greedy approach to action selection in reinforcement learning,   *1/1
which of the following values of ε makes the approach identical to a
purely greedy approach?

◉ $\varepsilon = 0$ ✓

○ $\varepsilon = 0.25$

○ $\varepsilon = 0.5$

○ $\varepsilon = 0.75$

○ $\varepsilon = 1$

Comments, if any

This form was created inside CS50.

Google Forms